Common scratch pool (Pool 00)

The *common scratch pool* is a pool that contains only scratch stacked volumes and serves as a reserve pool. For pools 01 - 32, you set it up to borrow scratch stacked cartridges from the common scratch pool (Pool 00) if a scratch shortage occurs. This setup can be done on a temporary or permanent basis.

Each pool can be defined to borrow a single media type (for example, JA, JB, JC, or JD), borrow mixed media, or have a first choice and a second choice. The borrowing options can be set by using the MI when you are defining stacked volume pool properties.

Note: The common scratch pool must have at least three scratch cartridges available to prevent low scratch warnings.

General-purpose pools (Pools 01 - 32)

A total of 32 general-purpose pools are available for each TS7700T cluster. These pools can contain empty, full, and filling stacked volumes. All physical volumes in a TS7700T cluster are distributed among available pools according to the physical volume range definitions within the MI.

The distribution is also based on the pools' borrow and return attribute settings. For example, you can have pools to be 01 - 32 be the target of a specific physical volume range to have a specific volume that is used for that pool. In this use case, the targeted pool should feature borrow/return disabled so that scratch volumes are not returned to the common scratch pool 00.

A pool's properties can be tailored individually by the administrator for various purposes. When initially creating these pools, it is important to ensure that the correct borrowing properties are defined for each one. For more information, see "Stacked volume pool properties" on page 51.

By default, one pool (Pool 01) is available and the TS7700T stores all newly created virtual volumes on any stacked volume that is available to it. This configuration creates an intermix of logical volumes from differing sources. The user cannot influence the physical location of the logical volumes within a physical pool, Defining more than one pool can be valuable for the following reasons:

- Need to separate different clients or LPAR data from each other.
- Need to segregate media types.
- Set up specific pools for Copy Export.
- Set up pool or pools for encryption with unique encryption key labels.
- Set a reclamation threshold at the pool level to match the reclamation needs of the data that is there.
- Set up reclamation cascading from one pool to another so that aged data ends up in pools with less reclamation activity.
- Set different numbers of physical back-end devices for premigration for different workloads.
- Assign specific physical volser ranges to specific workloads.