

to generate sample key-value pairs to use in your own benchmarks. See [libmemcached Command-Line Utilities](#) for details.

15.20.6.4 Controlling Transactional Behavior of the InnoDB memcached Plugin

Unlike traditional [memcached](#), the [daemon_memcached](#) plugin allows you to control durability of data values produced through calls to [add](#), [set](#), [incr](#), and so on. By default, data written through the [memcached](#) interface is stored to disk, and calls to [get](#) return the most recent value from disk. Although the default behavior does not offer the best possible raw performance, it is still fast compared to the SQL interface for [InnoDB](#) tables.

As you gain experience using the [daemon_memcached](#) plugin, you can consider relaxing durability settings for non-critical classes of data, at the risk of losing some updated values in the event of an outage, or returning data that is slightly out-of-date.

Frequency of Commits

One tradeoff between durability and raw performance is how frequently new and changed data is [committed](#). If data is critical, it should be committed immediately so that it is safe in case of an unexpected exit or outage. If data is less critical, such as counters that are reset after an unexpected exit or logging data that you can afford to lose, you might prefer higher raw throughput that is available with less frequent commits.

When a [memcached](#) operation inserts, updates, or deletes data in the underlying [InnoDB](#) table, the change might be committed to the [InnoDB](#) table instantly (if [daemon_memcached_w_batch_size=1](#)) or some time later (if the [daemon_memcached_w_batch_size](#) value is greater than 1). In either case, the change cannot be rolled back. If you increase the value of [daemon_memcached_w_batch_size](#) to avoid high I/O overhead during busy times, commits could become infrequent when the workload decreases. As a safety measure, a background thread automatically commits changes made through the [memcached](#) API at regular intervals. The interval is controlled by the [innodb_api_bk_commit_interval](#) configuration option, which has a default setting of 5 seconds.

When a [memcached](#) operation inserts or updates data in the underlying [InnoDB](#) table, the changed data is immediately visible to other [memcached](#) requests because the new value remains in the memory cache, even if it is not yet committed on the MySQL side.

Transaction Isolation

When a [memcached](#) operation such as [get](#) or [incr](#) causes a query or DML operation on the underlying [InnoDB](#) table, you can control whether the operation sees the very latest data written to the table, only data that has been committed, or other variations of transaction [isolation level](#). Use the [innodb_api_trx_level](#) configuration option to control this feature. The numeric values specified for this option correspond to isolation levels such as [REPEATABLE READ](#). See the description of the [innodb_api_trx_level](#) option for information about other settings.

A strict isolation level ensures that data you retrieve is not rolled back or changed suddenly causing subsequent queries to return different values. However, strict isolation levels require greater [locking](#) overhead, which can cause waits. For a NoSQL-style application that does not use long-running transactions, you can typically use the default isolation level or switch to a less strict isolation level.

Disabling Row Locks for memcached DML Operations

The [innodb_api_disable_rowlock](#) option can be used to disable row locks when [memcached](#) requests through the [daemon_memcached](#) plugin cause DML operations. By default, [innodb_api_disable_rowlock](#) is set to [OFF](#) which means that [memcached](#) requests row locks for [get](#) and [set](#) operations. When [innodb_api_disable_rowlock](#) is set to [ON](#), [memcached](#) requests a table lock instead of row locks.