## Characteristics for a zCX instance zFS file system

You also must allocate one zFS file system for each of your zCX instances. This is used to store the zCX appliance image and configuration data. In a sysplex environment, we recommend that you mount it under a mount point within the `/global` directory. In this way, you provide access to that file system from all participating systems. The required size for the file system is also >= 4GB.

## Planning for data set naming conventions

Naming conventions for the VSAM data sets are allocated during the execution of the `provision.xml` workflow, and these conventions require planning by you. Keep the following requirements in mind:

- The high-level qualifier (HLQ) of those data sets can contain up to 28 characters including the periods between the qualifiers.
- The `provision.xml` workflow allocates those data sets by using that HLQ in the format `HLQ.zcxinstancename.suffix`.
  - `zcxinstancename` is the jobname that you choose in the workflow, which you later use to start the zCX instance.
  - `suffix` is generated by the provisioning workflow and points to the function of the data set. The following data sets are provisioned (data set names are from our test environment):

    ```
    ZCX.REDB.ZCXRJ01.CONF
    ZCX.REDB.ZCXRJ01.DATA1
    ZCX.REDB.ZCXRJ01.DLOG1
    ZCX.REDB.ZCXRJ01.ROOT
    ZCX.REDB.ZCXRJ01.SWAP1
    ZCX.REDB.ZCXRJ01.ZFS
    ```
- If you plan to run different zCX instances (like production, development, or test) in a single z/OS system, you might want to distinguish the data sets by putting PROD, DEV, or TEST into the HLQ part of the data set names.

Figure 2-7 and Figure 2-8 on page 32, show you the structure of the data sets and their naming conventions.

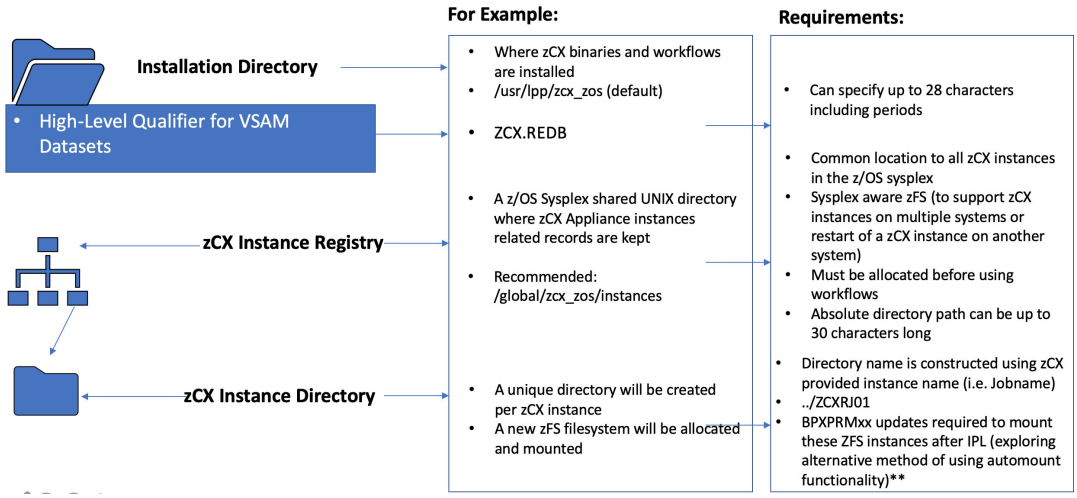


*Figure 2-7 zCX instance data set and zFS structure, part 1*