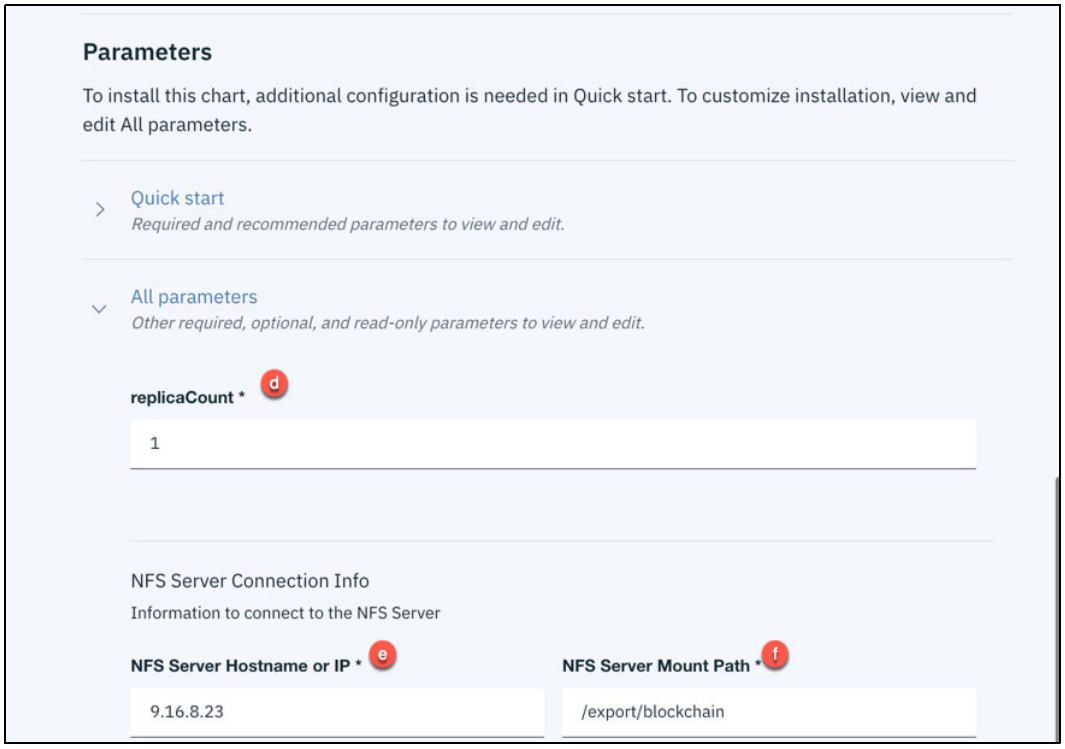

Figure 2-7 Setting up NFS server primer -3

The final values for callout letters g to k apply to Figure 2-8 on page 33.

Table 2-3 The final values, for callout letters g to k

g. `storageClassName` is the name of the StorageClass to create in Kubernetes by using the `nfs.io/nfs-client-provisioner`. (managed-nfs-storage)

h. `Make new Storage Class the default` makes the StorageClass that was created by the Helm chart into the default StorageClass in Kubernetes. This means that if a developer asks for a PersistentVolumeClaim and does not provide a StorageClass, the one created with storageClassName from callout letter **g** is used. It is good to have a default StorageClass to make PersistentVolume provisioning as easy as possible. However, if you have a different StorageClass that you want as default instead, uncheck this box. (checked)

i. `archiveOnDelete` signifies whether you want to archive PersistentVolumeClaims when they are deleted. (unchecked)

j. `Run as system-cluster-critical` is the recommended option that gives the provisioner priority over pods that rely on it. This prevents the problematic situation where the provisioner is evicted before the workloads that require it, which can lead to storage issues. (checked)
Note: If you select this option, remember to use kube-system as the `Target Namespace` value — shown by callout letter `b` in Figure 2-8 on page 33 — unless the deployment will not run.

k. `Resources` are set to recommended defaults that should work for most environments. If you need to change these values, remember that the limits cannot be lower than the requests.
Note: Limits and requests occur on a per container basis. This means that making multiple pods through changing the replicaCount of the helm release requires fulfilling the storage quota for each of those pods separately.