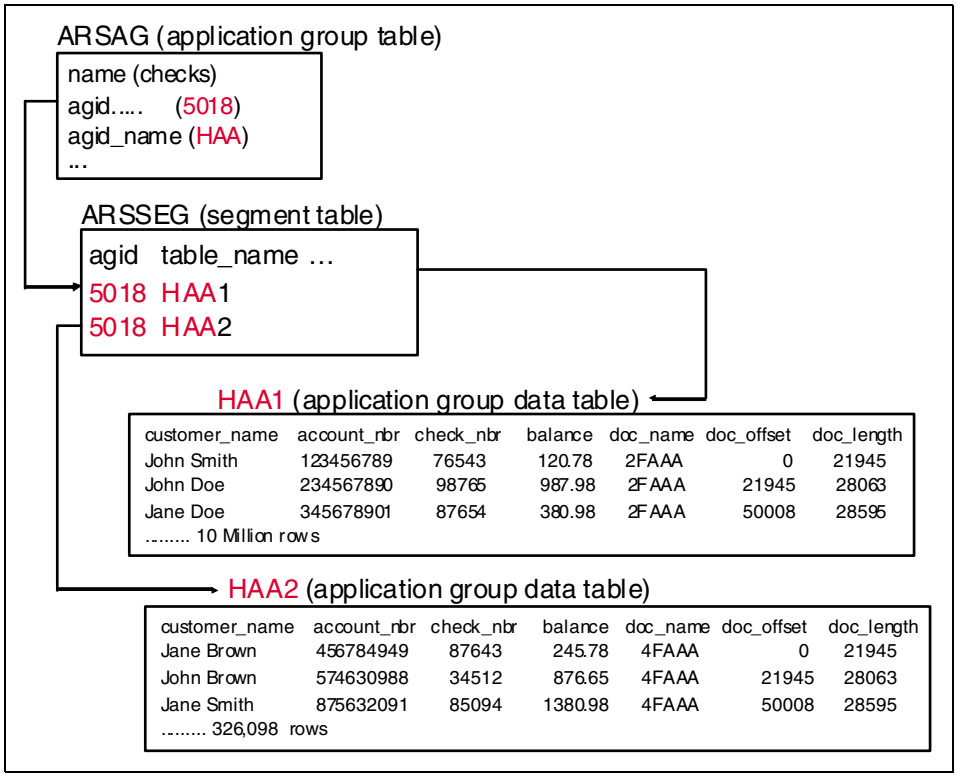

Figure 4-2 Relationship between system tables and data tables

This application group creates an application group data table every 10 million rows (based on the seg_rows value in the ARSAG table, which is not shown in Figure 4-2). During the data loading process, Content Manager OnDemand uses the agid and the agid_name to add a row into the segment table (ARSSEG) for every 10 million rows that are created in the application group data table. When the first data load occurs into the HAA application group, the index values for the stored documents are inserted into table HAA1. A row exists for each document that is loaded. When table HAA1 reaches its max_rows value (in this case 10 million rows), table HAA1 is closed and table HAA2 is opened.

The important pointer in the ARSSEG table is the name of the application group data table, *table_name*, where the index values (in this case, the four defined index values) are stored. The table_name consists of the agid_name from the ARSAG table, plus a counter.

Figure 4-2 shows the two rows that are created in the ARSSEG table: one row with the table_name HAA1 and another row with the table_name HAA2. Both HAA1 and HAA2 are the actual names of the application group data tables that are created.

The application group data table contains the *doc_name*, which is the actual container (storage object) for the individual document. The offset and the document length are also kept in this table. Figure 4-2 shows that the first row has an offset of 0, and the second row has an offset of the document length of the first row.

Figure 4-2 shows the relationship between the tables.

The architecture of relating one application group to one or more application group data tables allows Content Manager OnDemand an unlimited growth of index space. The maximum table size is a limitation of the database subsystem and must be configured for optimal performance.