

- For more information about the [CREATE TABLESPACE](#) and [ALTER TABLESPACE](#) statements, see [Section 13.1.21, “CREATE TABLESPACE Statement”](#), and [Section 13.1.10, “ALTER TABLESPACE Statement”](#).
3. Now it is possible to create a table whose unindexed columns are stored on disk using files in tablespace `ts_1`:

```
CREATE TABLE dt_1 (
  member_id INT UNSIGNED NOT NULL AUTO_INCREMENT PRIMARY KEY,
  last_name VARCHAR(50) NOT NULL,
  first_name VARCHAR(50) NOT NULL,
  dob DATE NOT NULL,
  joined DATE NOT NULL,
  INDEX(last_name, first_name)
)
TABLESPACE ts_1 STORAGE DISK
ENGINE NDBCLUSTER;
```

`TABLESPACE ts_1 STORAGE DISK` tells the NDB storage engine to use tablespace `ts_1` for data storage on disk.

Once table `ts_1` has been created as shown, you can perform [INSERT](#), [SELECT](#), [UPDATE](#), and [DELETE](#) statements on it just as you would with any other MySQL table.

It is also possible to specify whether an individual column is stored on disk or in memory by using a `STORAGE` clause as part of the column's definition in a [CREATE TABLE](#) or [ALTER TABLE](#) statement. `STORAGE DISK` causes the column to be stored on disk, and `STORAGE MEMORY` causes in-memory storage to be used. See [Section 13.1.20, “CREATE TABLE Statement”](#), for more information.

You can obtain information about the NDB disk data files and undo log files just created by querying the `FILES` table in the `INFORMATION_SCHEMA` database, as shown here:

```
mysql> SELECT
        FILE_NAME AS File, FILE_TYPE AS Type,
        TABLESPACE_NAME AS Tablespace, TABLE_NAME AS Name,
        LOGFILE_GROUP_NAME AS 'File group',
        FREE_EXTENTS AS Free, TOTAL_EXTENTS AS Total
        FROM INFORMATION_SCHEMA.FILES
        WHERE ENGINE='ndbcluster';
```

File	Type	Tablespace	Name	File group	Free	Total
./undo_1.log	UNDO LOG	lg_1	NULL	lg_1	0	4194304
./undo_2.log	UNDO LOG	lg_1	NULL	lg_1	0	3145728
./data_1.dat	DATAFILE	ts_1	NULL	lg_1	32	32
./data_2.dat	DATAFILE	ts_1	NULL	lg_1	48	48

```
4 rows in set (0.00 sec)
```

For more information and examples, see [Section 26.3.15, “The INFORMATION_SCHEMA FILES Table”](#).

Indexing of columns implicitly stored on disk. For table `dt_1` as defined in the example just shown, only the `dob` and `joined` columns are stored on disk. This is because there are indexes on the `id`, `last_name`, and `first_name` columns, and so data belonging to these columns is stored in RAM. Only nonindexed columns can be held on disk; indexes and indexed column data continue to be stored in memory. This tradeoff between the use of indexes and conservation of RAM is something you must keep in mind as you design Disk Data tables.

You cannot add an index to a column that has been explicitly declared `STORAGE DISK`, without first changing its storage type to `MEMORY`; any attempt to do so fails with an error. A column which *implicitly*